

MAJOR -2 PROJECT

END TERM REPORT

For

Designing a Scalable and Fault tolerant Web application

Submitted By

Specialization	SAP ID	Name
CCVT	500083440	Aahna Pandey
CCVT	500084089	Aarushi Arora
CCVT	500083149	Aditya Pilani



Department of Systemics

School Of Computer Science

UNIVERSITY OF PETROLEUM & ENERGY STUDIES,

DEHRADUN- 248007. Uttarakhand

Dr.Ambika Aggarwal
Project Guide

Dr. Neelu J. Ahuja
Cluster Head



School of Computer Science
University of Petroleum & Energy Studies, Dehradun

Synopsis Report

1. Project Title

Designing a Scalable and Fault tolerant Web application

2. Abstract

Web applications must provide outstanding performance and user experience, as well as high availability and resilience in the face of unforeseen difficulties, in order to survive in the current digital landscape. This calls for a fault-tolerant and scalable design methodology. A web application is said to be scalable if it can adapt to increases in user demand or data volume without experiencing a decline in performance. Contrarily, fault tolerance ensures that a system will continue to function even in the event of component failure or outside disturbances. In addition to meeting present needs, this presentation will cover future-proofing techniques and best practices for building web applications against growing loads and possible system failures by utilising AWS tools and resources. Through the strategic application of AWS's scalable resources and fault-tolerant architectures, developers can create web applications that not only withstand the test of time but also evolve with the technological landscape, ensuring that user demands are met with unwavering reliability and performance.

3. Introduction

- The success of a web application relies on its ability to remain operational and accessible, even when faced with unexpected challenges.
- Scalability and fault tolerance are pivotal concepts in achieving this goal.

- Scalability involves the application's capacity to adapt and perform well under increasing workloads.
- Fault tolerance ensures the application's availability even in the presence of failures or disruptions.
- Amazon Web Services (AWS) emerges as a cloud computing platform that offers a comprehensive suite of services tailored for contemporary web application development.
- The objective of this presentation is to explore how AWS resources and tools can be leveraged to design web applications that exhibit both scalability and fault tolerance.
- By effectively utilizing AWS services, developers can craft applications that seamlessly adjust to changing user demands and technological landscapes

4. Problem Statement

The difficulty lies in designing online applications that easily scale to user traffic while ensuring resilient availability in a digital environment marked by evolving user needs and technology uncertainties. Using the tactical capabilities of AWS services, this presentation discusses the requirement for designing such scalable and fault-tolerant web applications.

5. Objectives

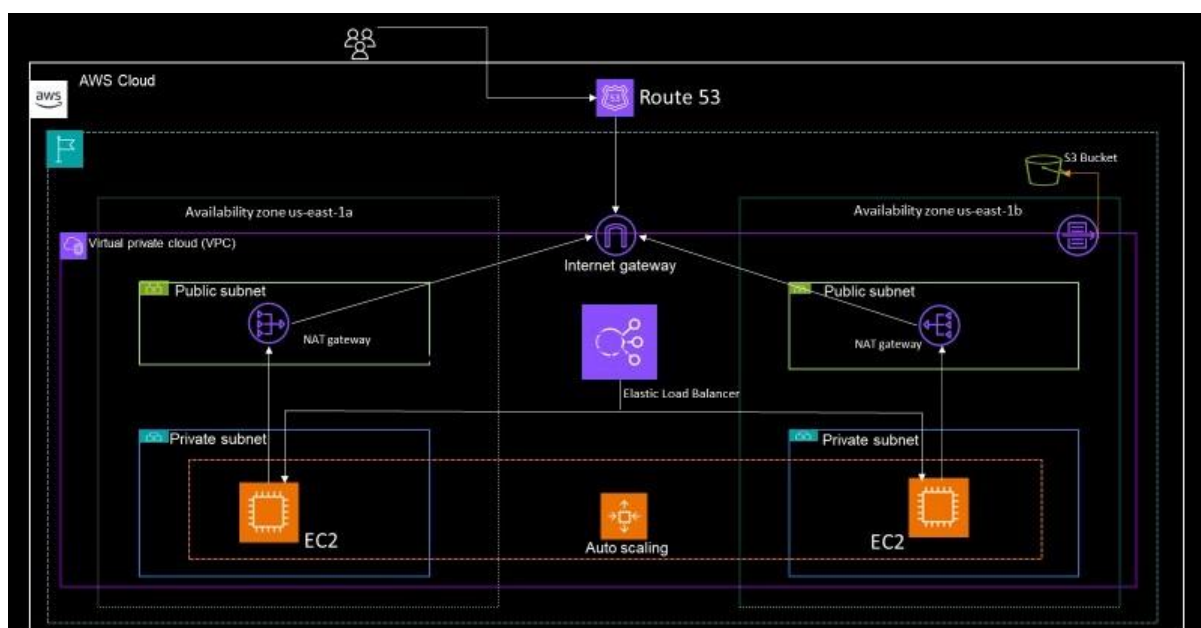
1. **Understand Principles:** Explore the fundamental concepts of scalability and fault tolerance in web application design using AWS services.
2. **AWS Tool Proficiency:** Gain proficiency in leveraging AWS services to create web applications that can scale dynamically and remain resilient.

3. **Architectural Strategies:** Discover effective architectural strategies for building applications capable of handling varying loads and ensuring fault tolerance.
4. **Real-World Application:** Learn through practical examples and case studies how to implement scalable and fault-tolerant design principles using AWS services.

6. Methodology

1. Create a VPC
2. Create public and private subnets
3. Create Route Table
4. Add Internet Gateway
5. Create EC2 instances
6. Create Target Groups
7. Create Elastic Load Balancer
8. Create Auto Scaling Groups

7. Flow Chart



8. RESULT:

Basic Settings

Security group name [Info](#)
WebServer SG
Name cannot be edited after creation.

Description [Info](#)
Allow access on ports 80 and 22

VPC [Info](#)
vpc-be482cc3

Inbound rules [Info](#)

Type Info	Protocol Info	Port range Info	Source Info	Description - optional Info	
HTTP	TCP	80	Anywhere...		Delete
SSH	TCP	22	My IP		Delete

Add rule

New EC2 Experience [X](#)

EC2 Dashboard
EC2 Global View
Events
Tags
Limits

Instances [New](#)
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances [New](#)
Dedicated Hosts
Scheduled Instances
Capacity Reservations

Images

Instances (1) [Info](#)

[Refresh](#) [Connect](#) [Instance state](#) [Actions](#) [Launch Instance](#)

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone
MyWebServer	i-0684c730c4c7bda2f	Running	t2.micro	2/2 checks passed	No alarms	us-east-1b

Select an instance

```
root@ip-172-31-8-1:/home/ec2-user
Using username "ec2-user".
Authenticating with public key "myec2key"

 _ _ | ( _ _ | )
 _ _ | ( _ _ | /
 _ _ | \ _ _ | _ |

Amazon Linux 2 AMI

https://aws.amazon.com/amazon-linux-2/
11 package(s) needed for security, out of 15 available
Run "sudo yum update" to apply all updates.
[ec2-user@ip-172-31-8-1 ~]$ sudo su
[root@ip-172-31-8-1 ec2-user]# yum update -y
```

New EC2 Experience

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Successfully terminated i-0684c730c4c7bda2f

Instances (1/1)

Search

Connect

Instance state

Actions

Launch Instance

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone
MyWebServer	i-0684c730c4c7bda2f	Shutting-down	t2.micro	2/2 checks passed	No alarms	us-east-1b

Instance: i-0684c730c4c7bda2f (MyWebServer)

Details

Security

Networking

Storage

Status checks

Monitoring

Tags

Instance summary

Instance ID

i-0684c730c4c7bda2f (MyWebServer)

IPv6 address

...

Hostname type

...

Public IPv4 address

3.235.161.207 [open address]

Instance state

Shutting-down

Private IP DNS name (IPv4 only)

...

Private IPv4 addresses

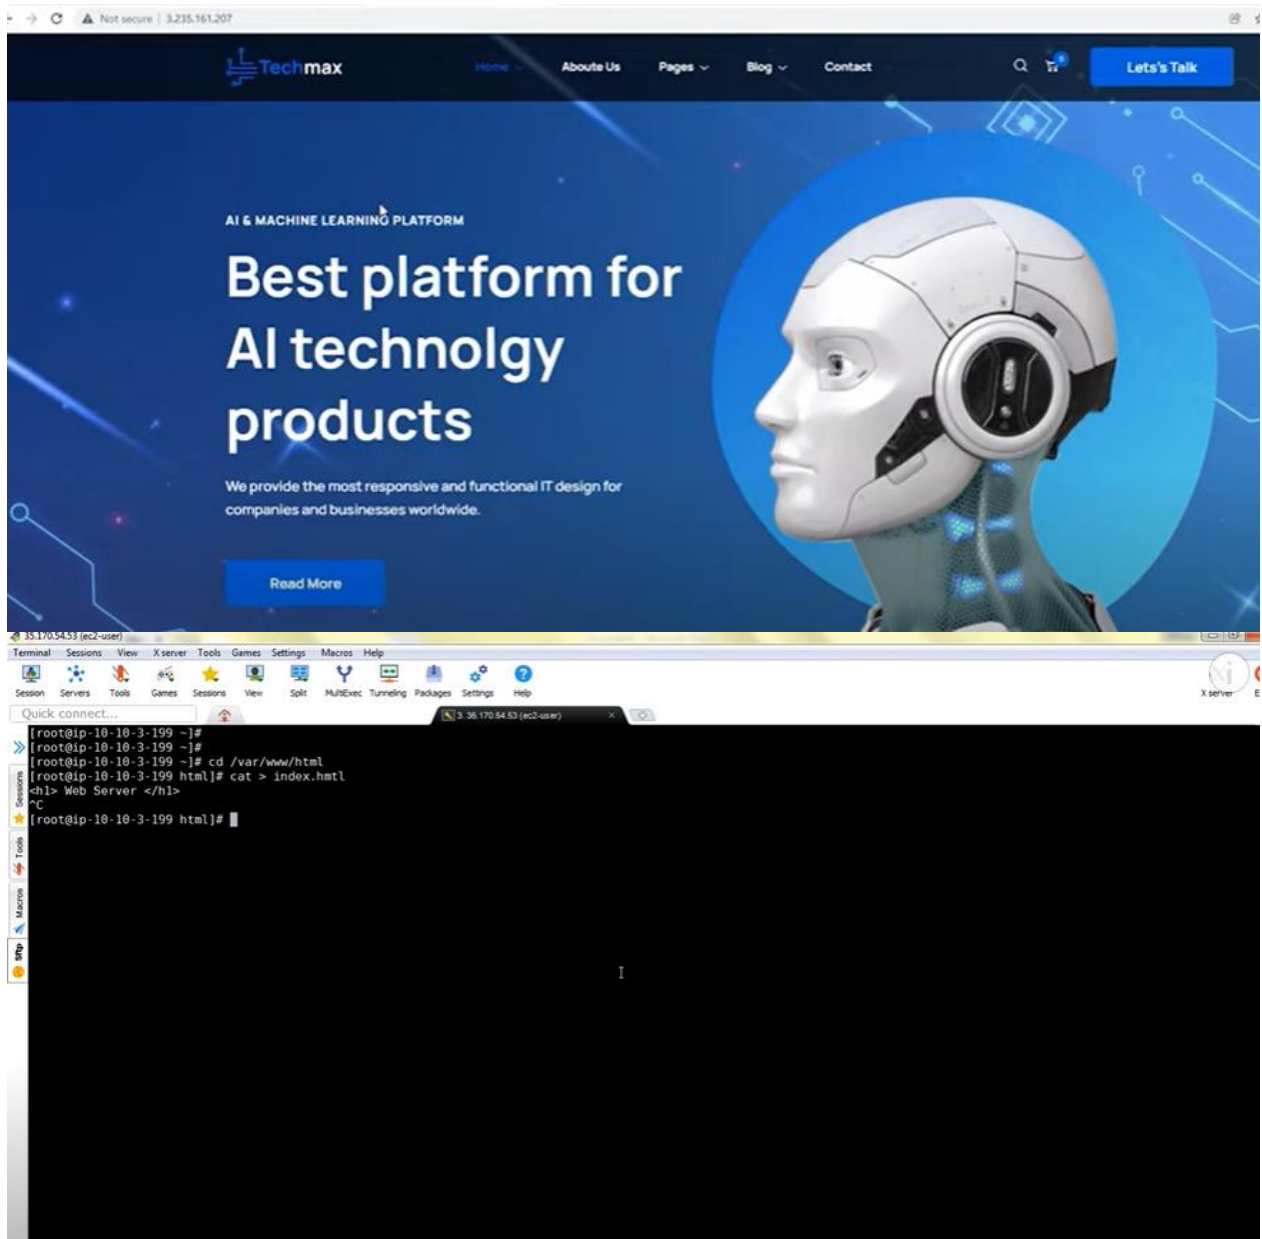
172.31.8.1

Public IPv4 DNS

ec2-3-235-161-207.compute-1.amazonaws.com [open address]




Answer private resource DNS name

...



Select load balancer type

Elastic Load Balancing supports three types of load balancers: Application Load Balancers, Network Load Balancers (new), and Classic Load Balancers. Choose the load balancer type that meets your needs. [Learn more](#) about which load balancer is right for you

Application Load Balancer	Network Load Balancer	Classic Load Balancer
		
Create	Create	Create
<p>Choose an Application Load Balancer when you need a flexible feature set for your web applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing, TLS termination and visibility features targeted at application architectures, including microservices and containers.</p> <p>Learn more ></p>	<p>Choose a Network Load Balancer when you need ultra-high performance and static IP addresses for your application. Operating at the connection level, Network Load Balancers are capable of handling millions of requests per second while maintaining ultra-low latencies.</p> <p>Learn more ></p>	<p>Choose a Classic Load Balancer when you have an existing application running in the EC2-Classical network.</p> <p>Learn more ></p>

Step 1: Define Load Balancer

Load Balancer Protocol	Load Balancer Port	Instance Protocol	Instance Port
HTTP	80	HTTP	80
<div>Add</div>			

Select Subnets

You will need to select a Subnet for each Availability Zone where you wish traffic to be routed by your load balancer. If you have instances in only one Availability Zone, please select at least two Subnets in different Availability Zones to provide higher availability for your load balancer.

VPC vpc-7af8ef02 (10.10.0.0/16) | DemoVPC

Available subnets

Actions	Availability Zone	Subnet ID	Subnet CIDR	Name
	us-east-1d	subnet-c76b02e8	10.10.4.0/24	Subnet-1D

Selected subnets

Actions	Availability Zone	Subnet ID	Subnet CIDR	Name
	us-east-1a	subnet-ba5e68f1	10.10.1.0/24	Subnet-1A
	us-east-1b	subnet-76274a25	10.10.2.0/24	Subnet-1B
	us-east-1c	subnet-03d9fc67	10.10.3.0/24	Subnet-1C

Step 2: Assign Security Groups

You have selected the option of having your Elastic Load Balancer inside of a VPC, which allows you to assign security groups to your load balancer. Please select the security groups to assign to this load balancer. This can be changed at any time.

Assign a security group: ☒ Create a new security group
☐ Select an existing security group

Security group name:
Description:

Type	Protocol	Port Range	Source
Custom TCP	TCP	80	Custom 0.0.0.0/0
<div>Add Rule</div>			

Step 4: Configure Health Check

Your load balancer will automatically perform health checks on your EC2 instances and only route traffic to instances that pass the health check. If an instance fails the health check, it is automatically removed from the load balancer. Customize the health check to meet your specific needs.

Ping Protocol:
Ping Port:
Ping Path:

Advanced Details

Response Timeout: seconds
Interval: seconds
Unhealthy threshold:
Healthy threshold:

1. Define Load Balancer 2. Assign Security Groups 3. Configure Security Settings 4. Configure Health Check 5. Add EC2 Instances 6. Add Tags 7. Review

Step 5: Add EC2 Instances

The table below lists all your running EC2 Instances. Check the boxes in the Select column to add those instances to this load balancer.

VPC vpc-7af8ef02 (10.10.0.0/16) | DemoVPC

Instance	Name	State	Security groups	Zone	Subnet ID	Subnet CIDR
<input checked="" type="checkbox"/>	i-0fa2b0503b62e97ba	running	WebSG	us-east-1c	subnet-03d9fc67	10.10.3.0/24

Availability Zone Distribution

1 instance in us-east-1c

☒ Enable Cross-Zone Load Balancing
☒ Enable Connection Draining seconds

Load Balancer Creation Status

✓

Successfully created load balancer

Load balancer **Demo-ELB** was successfully created.

Note: It may take a few minutes for your instances to become active in the new load balancer.

Close

Volumes

Snapshots

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

Owned by me

Filter by tags and attributes or search by keyword

1 to 1 of 1

Name	AMI Name	AMI ID	Source	Owner	Visibility	Status	Creation Date	Platform
Demo-Image		ami-0743717d	283396654523/...	283396654523	Private	available	January 25, 2018 at 10:52:5...	Other L

1. Choose AMI
2. Choose instance type
3. Configure details
4. Add storage
5. Configure security group
6. Review

Create Launch Configuration

Amazon EC2 provides a wide selection of instance types optimized to fit different use cases. Instances are virtual servers that can run applications. They have varying combinations of CPU, memory, storage, and network capacity, and give you the flexibility to choose the appropriate mix of resources for your applications. [Learn more](#) about instance types and how they can meet your computing needs.

Filter by:

All instance types

Current generation

Show/Hide Columns

Currently selected: t2.micro (Variable ECUs, 1 vCPUs, 2.5 GHz, Intel Xeon Family, 1 GiB memory, EBS only)

	Family	Type	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
<input type="checkbox"/>	General purpose	t2.nano	1	0.5	EBS only	-	Low to Moderate
<input checked="" type="checkbox"/>	General purpose	t2.micro Free tier eligible	1	1	EBS only	-	Low to Moderate
<input type="checkbox"/>	General purpose	t2.small	1	2	EBS only	-	Low to Moderate
<input type="checkbox"/>	General purpose	t2.medium	2	4	EBS only	-	Low to Moderate
<input type="checkbox"/>	General purpose	t2.large	2	8	EBS only	-	Low to Moderate
<input type="checkbox"/>	General purpose	t2.xlarge	4	16	EBS only	-	Moderate
<input type="checkbox"/>	General purpose	t2.2xlarge	8	32	EBS only	-	Moderate

1. Choose AMI 2. Choose Instance Type 3. Configure details 4. Add Storage 5. Configure Security Group 6. Review

Create Launch Configuration

Instance Type: t2.micro

Launch configuration details

Name: Demo-LC
Purchasing option: On demand
EBS Optimized: No
Monitoring: No
IAM role: None
Tenancy: Shared
Kernel ID: Use default
RAM Disk ID: Use default
User data:
IP Address Type: Only private

Storage

Security Groups

Select an existing key pair or create a new key pair

A key pair consists of a **public key** that AWS stores, and a **private key file** that you store. Together, they allow you to connect to your instance securely. For Windows AMIs, the private key file is required to obtain the password used to log into your instance. For Linux AMIs, the private key file allows you to securely SSH into your instance.

Note: The selected key pair will be added to the set of keys authorized for this instance. Learn more about [removing existing key pairs from a public AMI](#).

Choose an existing key pair

Select a key pair: WebKey

☒ I acknowledge that I have access to the selected private key file (WebKey pem), and that without this file, I won't be able to log into my instance.

Cancel Create launch configuration

1. Configure Auto Scaling group details 2. Configure scaling policies 3. Configure Notifications 4. Configure Tags 5. Review

Create Auto Scaling Group

None of the instances in this Auto Scaling group will be assigned a public IP address because you have not chosen to launch in your default VPC and subnet.

You can ensure a public IP address is assigned to instances launched with this configuration by selecting only default subnets of your default VPC.

[Learn more](#) about IP addressing in an Amazon VPC.

Advanced Details

Load Balancing ☒ Receive traffic from one or more load balancers [Learn about Elastic Load Balancing](#)

Classic Load Balancers: Demo-ELB

Target Groups:

Health Check Type: ☒ ELB ☐ EC2

Health Check Grace Period: 300 seconds

Monitoring: Amazon EC2 Detailed Monitoring metrics, which are provided at 1 minute frequency, are

1. Configure Auto Scaling group details 2. Configure scaling policies 3. Configure Notifications 4. Configure Tags 5. Review

Create Auto Scaling Group

You can optionally add scaling policies if you want to adjust the size (number of instances) of your group automatically. A scaling policy is a set of instructions for making such adjustments in response to an Amazon CloudWatch alarm that you assign to it. In each policy, you can choose to add or remove a specific number of instances or a percentage of the existing group size, or you can set the group to an exact size. When triggers, it will execute the policy and

☒ Keep this group at its size ☐ Use scaling policies to adjust the size

Scale between 2 and

Increase Group Size

Name:
Execute policy when:
Take the action:
Instances need:
Create a simple scaling policy

Create Alarm

You can use CloudWatch alarms to be notified automatically whenever metric data reaches a level you define.

To edit an alarm, first choose whom to notify and then define when the notification should be sent.

☒ Send a notification to: No SNS topics found

Whenever: Average of CPU Utilization

Is: >= Percent


For at least: 1 consecutive period(s) of 5 Minutes

Name of alarm: awssec2-Demo-ASG-CPU-Utilization

CPU Utilization Percent

Cancel Create Alarm

Auto Scaling group creation status



Initiating Auto Scaling group creation
Please do not close your browser

Creating Auto Scaling group... Successful

Retrieving Auto Scaling group details... Successful

Creating Scaling policies...

EC2 Dashboard

Instances

Launch Templates

Spot Requests

Reserved Instances

Dedicated Hosts

Scheduled Instances

Images

AMIs

Bundle Tasks

Elastic Block Store

Volumes

Snapshots

Network & Security

Security Groups

Create Auto Scaling group

Actions

Filter:

1 to 1 of 1 Auto Scaling Groups

Name	Launch Configuration /	Instances	Desired	Min	Max	Availability Zones	Default Cooldown	Health Check Grace
Demo-ASG	Demo-LC	2	2	2	5	us-east-1a, us-east-1b, us-e...	300	300

Auto Scaling Group: Demo-ASG

Details

Activity History

Scaling Policies

Instances

Monitoring

Notifications

Tags

Scheduled Actions

Lifecycle Hooks

Launch Configuration

Demo-LC

Launch Template

Launch Template Version

Load Balancers

Demo-ELB

Target Groups

Desired

2

Availability Zone(s)

us-east-1a, us-east-1b, us-east-1c, us-east-1d

EC2 Dashboard

Events

Tags

Reports

Limits

INSTANCES

Instances

Launch Templates

Spot Requests

Reserved Instances

Dedicated Hosts

Scheduled Instances

IMAGES

AMIs

Bundle Tasks

ELASTIC BLOCK STORE

Volumes

Snapshots

NETWORK & SECURITY

Create Load Balancer

Actions

Filter:

1 to 1 of 1

Name	DNS name	State	VPC ID	Availability Zones	Type
Demo-ELB	Demo-ELB-1757347503 us-...		vpc-7af9ef02	us-east-1a, us-east-1b,...	classic

Connection Draining: Enabled, 300 seconds (Edit)

Edit Instances

Instance ID	Name	Availability Zone	Status	Actions
i-074497cc2dc96277c	Demo-ASG	us-east-1b	OutOfService ⓘ	Remove from Load Balancer
i-0fb1c1716346d726	Demo-ASG	us-east-1c	InService ⓘ	Remove from Load Balancer
i-0fa2b0503b62e97ba	WebServer	us-east-1c	InService ⓘ	Remove from Load Balancer

Edit Availability Zones

9. References

- <https://cloudknight.medium.com/building-a-fault-tolerant-web-application-on-aws-a-step-by-step-guide-to-using-cloudformation-590988a99eb8>
- <https://aws.plainenglish.io/create-a-scalable-resilient-fault-tolerant-3-tier-architecture-9cdb3064b2de>
- <https://awstip.com/designing-a-fault-tolerant-web-application-building-a-highly-available-3-tier-architecture-e740765dd069>